Abstract : Real-time Log Analysis Using Hadoop and Spark

The **Real-time Log Analysis** project is a Big Data solution designed to collect, process, and analyze system and application logs in real time using the Hadoop and Apache Spark ecosystem. It enables organizations to monitor application performance, detect anomalies or failures, and gain insights into user behavior or system usage patterns — all in near real time for proactive decision-making.[1]

Key Features:

- Collect streaming log data from servers or applications using tools like **Flume** or **Kafka**.[2]

- Process logs in real time using **Apache Spark Streaming**.[3]

- Perform parsing, filtering, and aggregation of log records.
- Detect anomalies, frequent errors, or traffic spikes on the fly.
- Store processed logs in HDFS for batch analysis or historical trend tracking.
- Visualize trends, alerts, and system health using tools like **Grafana**, **Kibana**, or **Power BI**.[4]

Tools & Technologies Used:

- **Apache Kafka** – for ingesting real-time logs[5]

- **Apache Flume** – (alternative to Kafka) for log collection
- **Apache Spark Streaming** – for real-time processing[6]

- **HDFS** – for storage of raw and processed logs
- **YARN** – for cluster resource management[7]

- **Hive/Spark SQL** – for querying log data
- **Elasticsearch + Kibana** – (optional) for search and visualization
- **Scala or Python** – for Spark programming

Input Data:

- Web server logs (e.g., Apache/Nginx logs: IP address, timestamp, URL, response code)[8]

- Application logs (e.g., error logs, transaction logs)

- System logs (e.g., CPU/memory usage, disk I/O logs)
- Streaming data from log-generating services or APIs

Output/Reports:

- Real-time error and event dashboards
- Aggregated reports: most visited URLs, error codes, user IPs
- Alerting on anomalies (e.g., traffic spikes, error surges)
- Hourly/daily log summaries stored in Hive/Spark tables
- Trend graphs and logs heatmaps for performance monitoring

Processing Logic:

- **Spark Streaming job** runs in a loop consuming batches of logs from Kafka every few seconds.
- Apply **filter conditions** (e.g., only ERROR or 500 status logs).
- Use **window operations** for time-based aggregations (e.g., every 10 mins).
- Store results in **HDFS** or **Elasticsearch**.

24M11MC243

Kamadi Pradeep Kumar

Aditya University